

Assignment 3: Indexing EotW

Group 10 – Emerald

Kanchan Chowdhari

Payal Kaur

Purva Khandelwal

Samiksha Mhatre

1. Idea 1: Investigating Selectivity

To demonstrate this experiment, we have queried upon one of our Group project's table- PATIENT_RECORDS.

1. Let's **display the Patient name, Patient ID, gender and patient status** from Patient_Records table **without indexing**.

```
SELECT patient_id, pfname, pname,  
p_gender, patient_status  
FROM patient_records  
WHERE p_gender IN ('M', 'F') AND patient_status IN ('Active', 'Inactive');
```

The screenshot displays a database query execution interface. At the top, the SQL query is entered in a text area:

```
SELECT patient_id, pfname, pname  
p_gender, patient_status  
FROM patient_records  
WHERE p_gender IN ('M', 'F') AND patient_status IN ('Active', 'Inactive');
```

Below the query, the execution plan is shown. The top part of the plan is a table with the following columns: OPERATION, OBJECT_NAME, OPTIONS, CARDINALITY, and COST.

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				2048 30
TABLE ACCESS	PATIENT_RECORDS	FULL		2048 30

Below the table, the execution plan is visualized as a tree structure. The root node is 'SELECT STATEMENT', which branches into 'TABLE ACCESS' and 'Filter Predicates'. The 'Filter Predicates' node branches into 'AND', which further branches into 'OR' and 'OR'. The 'OR' nodes branch into 'PATIENT_STATUS='Active'' and 'PATIENT_STATUS='Inactive'', and 'P_GENDER='F'' and 'P_GENDER='M'' respectively. The 'TABLE ACCESS' node branches into 'Other XML', which further branches into '{info}' and 'info type="db_version"', and 'info type="parse_schema"'.

At the bottom of the interface, the status bar shows 'Line 14 Column 21 | Insert | Modified | Windows: C'.

2. To check how bitmap indexing affects the performance of the query, we **have created indexes on the gender and patient status columns**.

```
CREATE BITMAP INDEX pat_index  
ON patient_records (p_gender);
```

```
CREATE BITMAP INDEX status_index  
ON patient_records (patient_status);
```

- Now we get the explain plan for the below query wherein we have **eliminated the columns like Patient first name, Patient last name**, to get the indexing applied on indexed fields.

Script Output x Query Result x Explain Plan x Query Result 1 x							
All Rows Fetched: 32 in 0.071 seconds							
PLAN_TABLE_OUTPUT							
1	Plan hash value: 311294621						
2							
3	-----						
4	Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
5	-----						
6	0	SELECT STATEMENT		2048	30720	15 (0)	00:00:01
7	* 1	VIEW	index\$ join\$ 001	2048	30720	15 (0)	00:00:01
8	* 2	HASH JOIN					
9	* 3	HASH JOIN					
10	4	INLIST ITERATOR					
11	5	BITMAP CONVERSION TO ROWIDS		2048	30720	2 (0)	00:00:01
12	* 6	BITMAP INDEX SINGLE VALUE	PSTATUS_INDX				
13	7	INLIST ITERATOR					
14	8	BITMAP CONVERSION TO ROWIDS		2048	30720	2 (0)	00:00:01
15	* 9	BITMAP INDEX SINGLE VALUE	PAT_INDEX				
16	10	INDEX FAST FULL SCAN	PATIENT_RECORDS_PK	2048	30720	11 (0)	00:00:01
17	-----						
18							

Conclusion-When there were no indexes, there was full access to the table for all the selected fields. After adding the bit map indexes for low cardinality columns like gender, status, we could only select these indexed fields to be able to implement bitmap indexing. Hence by shrinking the selective lookup to indexed fields thereby increased the performance by reducing the query cost from 30 to 15.

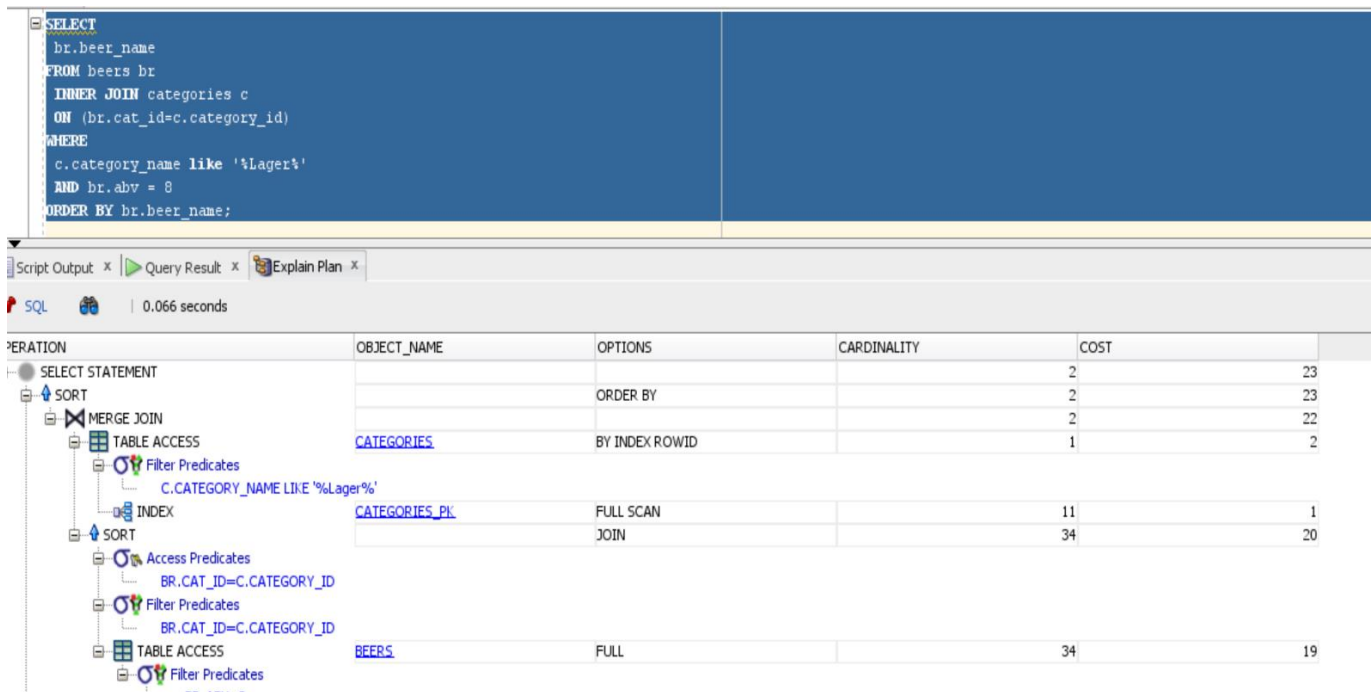
2. Idea 2: Start Simple and show that Indexing works

Let's take two examples/queries to demonstrate this.

1. **Display all beer names that belong to a category with a name containing "Lager" somewhere in the name and have an alcohol by volume (ABV) of eight. Show the beer names in alphabetical order.**

```
SELECT
    br.beer_name
FROM
    beers br
    INNER JOIN categories c
        ON (br.cat_id=c.category_id)
WHERE c.category_name LIKE '%Lager%'
AND br.abv = 8
ORDER BY
    br.beer_name;
```

The explain plan for this query is as shown in below:



OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				23
SORT		ORDER BY		23
MERGE JOIN				22
TABLE ACCESS	CATEGORIES	BY INDEX ROWID		2
Filter Predicates				
		C.CATEGORY_NAME LIKE '%Lager%'		
INDEX	CATEGORIES_PK	FULL SCAN		11
SORT		JOIN		34
Access Predicates				
		BR.CAT_ID=C.CATEGORY_ID		
Filter Predicates				
		BR.CAT_ID=C.CATEGORY_ID		
TABLE ACCESS	BEERS	FULL		34
Filter Predicates				
		BR.ABV = 8		

As seen from the plan we have full access to tables. In order to remove them we **can create index and check the explain plan again.**

- ```

SELECT
 b.beer_name,
 COUNT(bm.award_id) as TOTAL_AWARDS
FROM
 beer_award_mapping bm
 INNER JOIN beers b
 ON (bm.beer_id=b.beer_id)
 INNER JOIN awards a
 ON (bm.award_id=a.award_id)
 INNER JOIN competitions c
 ON (bm.competition_id=c.competition_id)
WHERE c.year = '2014'
AND award_level IN ('Gold','Silver','Bronze')
GROUP BY
 b.beer_name
ORDER BY
 TOTAL_AWARDS desc;

```

```

1 PLAN_TABLE_OUTPUT
2
3 Plan hash value: 2532414249
4
5 -----
6 | Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time |
7 -----
8 | 0 | SELECT STATEMENT | | 16 | 1104 | 27 (12) | 00:00:01 |
9 | 1 | SORT ORDER BY | | 16 | 1104 | 27 (12) | 00:00:01 |
10 | 2 | HASH GROUP BY | | 16 | 1104 | 27 (12) | 00:00:01 |
11 | 3 | NESTED LOOPS | | 16 | 1104 | 25 (4) | 00:00:01 |
12 | 4 | NESTED LOOPS | | 16 | 1104 | 25 (4) | 00:00:01 |
13 |* 5 | HASH JOIN | | 16 | 496 | 9 (12) | 00:00:01 |
14 | 6 | MERGE JOIN | | 17 | 357 | 6 (17) | 00:00:01 |
15 |* 7 | TABLE ACCESS BY INDEX ROWID | COMPETITONS | 2 | 18 | 2 (0) | 00:00:01 |
16 | 8 | INDEX FULL SCAN | COMPETITIONS_PK | 17 | | 1 (0) | 00:00:01 |
17 |* 9 | SORT JOIN | | 143 | 1716 | 4 (25) | 00:00:01 |
18 | 10 | TABLE ACCESS FULL | BEER_AWARD_MAPPING | 143 | 1716 | 3 (0) | 00:00:01 |
19 |* 11 | TABLE ACCESS FULL | AWARDS | 20 | 200 | 3 (0) | 00:00:01 |
20 |* 12 | INDEX UNIQUE SCAN | BEERS_PK | 1 | | 0 (0) | 00:00:01 |
21 | 13 | TABLE ACCESS BY INDEX ROWID | BEERS | 1 | 38 | 1 (0) | 00:00:01 |
22 -----
23
24 Predicate Information (identified by operation id):
25 -----
26 5 - access("BM"."AWARD_ID"="A"."AWARD_ID")
27 7 - filter("C"."YEAR"='2014')
28 9 - access("BM"."COMPETITION_ID"="C"."COMPETITION_ID")
29 filter("BM"."COMPETITION_ID"="C"."COMPETITION_ID")
30 11 - filter("A"."AWARD_LEVEL"='Bronze' OR "A"."AWARD_LEVEL"='Gold' OR
31 "A"."AWARD_LEVEL"='Silver')
32 12 - access("BM"."BEER_ID"="B"."BEER_ID")

```

As seen from the plan we have full access to tables. In order to remove them we **can create index and check the explain plan again.**

The following indexes were created on the table columns

```
CREATE INDEX bam_compid_idx
ON beer_award_mapping(competition_id);
```

```
CREATE INDEX aw_awardlevel_idx
ON awards(award_level);
```

```
CREATE INDEX comp_year_idx
ON competitions(year);
```

| PLAN_TABLE_OUTPUT |                                                                                                   |                                     |                    |      |       |             |          |
|-------------------|---------------------------------------------------------------------------------------------------|-------------------------------------|--------------------|------|-------|-------------|----------|
| 3                 | -----                                                                                             |                                     |                    |      |       |             |          |
| 4                 | Id                                                                                                | Operation                           | Name               | Rows | Bytes | Cost (%CPU) | Time     |
| 5                 | -----                                                                                             |                                     |                    |      |       |             |          |
| 6                 | 0                                                                                                 | SELECT STATEMENT                    |                    | 16   | 1104  | 24 (9)      | 00:00:01 |
| 7                 | 1                                                                                                 | SORT ORDER BY                       |                    | 16   | 1104  | 24 (9)      | 00:00:01 |
| 8                 | 2                                                                                                 | HASH GROUP BY                       |                    | 16   | 1104  | 24 (9)      | 00:00:01 |
| 9                 | 3                                                                                                 | NESTED LOOPS                        |                    | 16   | 1104  | 22 (0)      | 00:00:01 |
| 10                | 4                                                                                                 | NESTED LOOPS                        |                    | 16   | 1104  | 22 (0)      | 00:00:01 |
| 11                | * 5                                                                                               | HASH JOIN                           |                    | 16   | 496   | 6 (0)       | 00:00:01 |
| 12                | 6                                                                                                 | NESTED LOOPS                        |                    | 17   | 357   | 4 (0)       | 00:00:01 |
| 13                | 7                                                                                                 | NESTED LOOPS                        |                    | 17   | 357   | 4 (0)       | 00:00:01 |
| 14                | 8                                                                                                 | TABLE ACCESS BY INDEX ROWID BATCHED | COMPETITONS        | 2    | 18    | 2 (0)       | 00:00:01 |
| 15                | * 9                                                                                               | INDEX RANGE SCAN                    | COMP_YEAR_IDX      | 2    |       | 1 (0)       | 00:00:01 |
| 16                | * 10                                                                                              | INDEX RANGE SCAN                    | BAM_COMPID_IDX     | 8    |       | 0 (0)       | 00:00:01 |
| 17                | 11                                                                                                | TABLE ACCESS BY INDEX ROWID         | BEER_AWARD_MAPPING | 8    | 96    | 1 (0)       | 00:00:01 |
| 18                | 12                                                                                                | INLIST ITERATOR                     |                    |      |       |             |          |
| 19                | 13                                                                                                | TABLE ACCESS BY INDEX ROWID BATCHED | AWARDS             | 20   | 200   | 2 (0)       | 00:00:01 |
| 20                | * 14                                                                                              | INDEX RANGE SCAN                    | AW_AWARDLEVEL_IDX  | 20   |       | 1 (0)       | 00:00:01 |
| 21                | * 15                                                                                              | INDEX UNIQUE SCAN                   | BEERS_PK           | 1    |       | 0 (0)       | 00:00:01 |
| 22                | 16                                                                                                | TABLE ACCESS BY INDEX ROWID         | BEERS              | 1    | 38    | 1 (0)       | 00:00:01 |
| 23                | -----                                                                                             |                                     |                    |      |       |             |          |
| 24                |                                                                                                   |                                     |                    |      |       |             |          |
| 25                | Predicate Information (identified by operation id):                                               |                                     |                    |      |       |             |          |
| 26                | -----                                                                                             |                                     |                    |      |       |             |          |
| 27                |                                                                                                   |                                     |                    |      |       |             |          |
| 28                | 5 - access("BM"."AWARD_ID"="A"."AWARD_ID")                                                        |                                     |                    |      |       |             |          |
| 29                | 9 - access("C"."YEAR"='2014')                                                                     |                                     |                    |      |       |             |          |
| 30                | 10 - access("BM"."COMPETITION_ID"="C"."COMPETITION_ID")                                           |                                     |                    |      |       |             |          |
| 31                | 14 - access("A"."AWARD_LEVEL"='Bronze' OR "A"."AWARD_LEVEL"='Gold' OR "A"."AWARD_LEVEL"='Silver') |                                     |                    |      |       |             |          |
| 32                | 15 - access("BM"."BEER_ID"="B"."BEER_ID")                                                         |                                     |                    |      |       |             |          |
| --                |                                                                                                   |                                     |                    |      |       |             |          |

After creating the indexes and checking the explain plan, it **shows full access on table is removed and the cost is reduced and the query is optimized for faster retrieval of records.**

**Conclusion:** When there were no indexes on the table columns, full access took place until it fetched the required data. Since these tables have large data sets, selection on full tables impact the performance. Adding a B-tree index on these table columns eliminates the full table access and reduces the cost of the query resulting in faster retrieval of records.

### 3. Idea 6: Database programming

A stored procedure is a group of SQL statements that can be created and stored in database. It can be shared by multiple programs and can be reused. Stored procedures may return result sets and can be processed using cursors, they may contain variables for processing data, they may be also used for data generation or manipulation.

For demonstrating this,

1. Let's create a simple procedure that will generate data for a table fetching few values from other tables using a cursor.

```
create or replace PROCEDURE PROC_USER_ROLE_MAP_INSERT AS

CURSOR cur_user IS
SELECT user_id FROM user_master
WHERE user_id <> 100
ORDER BY user_id ASC;

CURSOR cur_patient IS
SELECT patient_id FROM patient_records
ORDER BY patient_id ASC;

CURSOR cur_vendor IS
SELECT vendor_id FROM vendor_details
ORDER BY vendor_id ASC;

BEGIN
FOR rec_user IN cur_user
LOOP
 INSERT INTO user_role_mapping (role_id,user_id) VALUES (2,rec_user.user_id);
END LOOP;
COMMIT;
dbms_output.put_line('Record inserted successfully for Users');

FOR rec_pat IN cur_patient
LOOP
 INSERT INTO user_role_mapping (role_id,user_id) VALUES (3,rec_pat.patient_id);
END LOOP;
COMMIT;
dbms_output.put_line('Record inserted successfully for Patients');

FOR rec_vendor IN cur_vendor
LOOP
 INSERT INTO user_role_mapping (role_id,user_id) VALUES (4,rec_vendor.vendor_id);
END LOOP;
COMMIT;
dbms_output.put_line('Record inserted successfully for Vendors');

EXCEPTION
 WHEN OTHERS THEN
 dbms_output.put_line('THIS IS EXCEPTION SECTION');
END PROC_USER_ROLE_MAP_INSERT;
```



After executing the stored procedure **PROC\_USER\_ROLE\_MAP\_INSERT**, the above data is generated for **USER\_ROLE\_MAPPING** table by fetching few details from **USER\_MASTER**, **PATIENT\_RECORDS** and **VENDOR\_DETAILS** tables using **cursors**.

```
select * from user_role_mapping order by role_id,user_id asc
```

| Query Result x                         |         |         |
|----------------------------------------|---------|---------|
| SQL   Fetched 50 rows in 0.036 seconds |         |         |
|                                        | ROLE_ID | USER_ID |
| 1                                      | 1       | 100     |
| 2                                      | 2       | 101     |
| 3                                      | 2       | 102     |
| 4                                      | 2       | 103     |
| 5                                      | 2       | 104     |
| 6                                      | 2       | 105     |
| 7                                      | 2       | 106     |
| 8                                      | 2       | 107     |
| 9                                      | 2       | 108     |
| 10                                     | 2       | 109     |
| 11                                     | 2       | 110     |
| 12                                     | 2       | 111     |
| 13                                     | 2       | 112     |
| 14                                     | 2       | 113     |
| 15                                     | 2       | 114     |
| 16                                     | 2       | 115     |
| 17                                     | 3       | 1000    |
| 18                                     | 3       | 1001    |
| 19                                     | 3       | 1002    |
| 20                                     | 3       | 1003    |
| 21                                     | 3       | 1004    |
| 22                                     | 3       | 1005    |

After executing the select query from user\_role\_mapping table, we can see records in the table which are inserted by the procedure.

2. Let's create a simple procedure that will generate data for a table and after generating records will also update few columns of the same table.

```
create or replace PROCEDURE proc_data_generation
IS
BEGIN
 dbms_output.put_line('This is begin section');
 FOR I IN 1..200
 LOOP
 --insert records into orders table
 INSERT INTO ORDERS (ORDER_ID
 ,ORDER_DATE
 ,DELIVERY_DUE_DATE
 ,PAYMENT_MODE
 ,TOTAL_PRICE
 ,USER_ID)
 VALUES (ORDER_SEQ.NEXTVAL
 ,SYSDATE
 ,SYSDATE
 , 'ONLINE'
 ,0
 ,5);

 COMMIT;
 dbms_output.put_line('RECORD INSERTED SUCCESSFULLY');

 --manipulating records in orders table
 UPDATE ORDERS
 SET DELIVERY_DUE_DATE = '01-JUN-21';

 UPDATE ORDERS
 SET TOTAL_PRICE = ROUND(DBMS_RANDOM.VALUE(1,80000));

 COMMIT;
 END LOOP;

EXCEPTION
 WHEN OTHERS THEN
 dbms_output.put_line('THIS IS EXCEPTION SECTION');
END;
```

After executing the stored procedure **PROC\_DATA\_GENERATION**, data is generated for **ORDERS** table. Once the default data is inserted into the table, we manipulate the two columns Delivery\_due\_Date and Total\_Price by writing the update statement inside the stored procedure.

| select * from orders order by order_id asc |          |            |                   |              |             |         |
|--------------------------------------------|----------|------------|-------------------|--------------|-------------|---------|
| Query Result x                             |          |            |                   |              |             |         |
| SQL   Fetched 50 rows in 0.033 seconds     |          |            |                   |              |             |         |
|                                            | ORDER_ID | ORDER_DATE | DELIVERY_DUE_DATE | PAYMENT_MODE | TOTAL_PRICE | USER_ID |
| 1                                          | 2        | 06-JAN-17  | 01-JUN-21         | ONLINE       | 40310       | 5       |
| 2                                          | 3        | 26-AUG-01  | 01-JUN-21         | ONLINE       | 58260       | 5       |
| 3                                          | 4        | 22-NOV-20  | 01-JUN-21         | ONLINE       | 9272        | 5       |
| 4                                          | 5        | 02-AUG-09  | 01-JUN-21         | ONLINE       | 76450       | 5       |
| 5                                          | 6        | 05-FEB-21  | 01-JUN-21         | ONLINE       | 49500       | 5       |
| 6                                          | 7        | 12-DEC-19  | 01-JUN-21         | ONLINE       | 15540       | 5       |
| 7                                          | 8        | 13-OCT-04  | 01-JUN-21         | ONLINE       | 5832        | 5       |
| 8                                          | 9        | 21-DEC-20  | 01-JUN-21         | ONLINE       | 73570       | 5       |
| 9                                          | 10       | 01-APR-07  | 01-JUN-21         | ONLINE       | 67380       | 5       |
| 10                                         | 11       | 20-OCT-09  | 01-JUN-21         | ONLINE       | 38420       | 5       |
| 11                                         | 12       | 06-OCT-09  | 01-JUN-21         | ONLINE       | 50180       | 5       |

After executing the select query from orders table we can see records in the table which are inserted and updated by the procedure.

**Conclusion:** Stored procedures can be created to perform various operations like data generation, manipulation, computation, and many more operations in the database that can be reused and shared by multiple programs/applications.